

Assignment 6

Vishal Batvia

CMPT 762 Computer Vision

Author Note

Kaggle Name: - VishalBatavia

Assignment 6

1. Part 1 BaseNet on CIFAR 100

1.1. Improved Network Architecture

```

BaseNet(
  (conv_layer): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (9): ReLU(inplace=True)
    (10): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (12): ReLU(inplace=True)
    (13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (16): ReLU(inplace=True)
    (17): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (19): ReLU(inplace=True)
    (20): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (22): ReLU(inplace=True)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (24): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (25): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (26): ReLU(inplace=True)
    (27): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (28): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (29): ReLU(inplace=True)
    (30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (31): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (32): ReLU(inplace=True)
    (33): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (34): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (35): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (36): ReLU(inplace=True)
    (37): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (38): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (39): ReLU(inplace=True)
    (40): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (41): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (42): ReLU(inplace=True)
    (43): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (fc_layer): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=1024, out_features=512, bias=True)
    (2): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (3): ReLU(inplace=True)
    (4): Dropout(p=0.5, inplace=False)
    (5): Linear(in_features=512, out_features=512, bias=True)

    (6): BatchNorm1d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): ReLU(inplace=True)
    (8): Linear(in_features=512, out_features=100, bias=True)
  )
)

```

1.2. Improvements in accuracy

The accuracy with the vanilla model Basenet was 21% on the test set on Kaggle while the accuracy of the final improved network was 64.5% on test set on Kaggle and 62% on validation set while training.

```
Accuracy of the network on the val images: 61 %  
[146] loss: 0.997  
Accuracy of the network on the val images: 60 %  
[147] loss: 0.985  
Accuracy of the network on the val images: 62 %  
[148] loss: 1.001  
Accuracy of the network on the val images: 61 %  
[149] loss: 0.988  
Accuracy of the network on the val images: 61 %  
[150] loss: 0.999  
Accuracy of the network on the val images: 60 %  
Finished Training
```

1.3. Changes in network

1.3.1. Data Augmentation:

```
train_transform = transforms.Compose([transforms.RandomCrop(32, padding=4),  
                                     transforms.RandomHorizontalFlip(),  
                                     transforms.ToTensor(),  
                                     transforms.Normalize([0.4914, 0.4822, 0.4465],  
                                                         [0.2023, 0.1994, 0.2010]))  
test_transform = transforms.Compose([transforms.ToTensor(),  
                                    transforms.Normalize([0.4914, 0.4822, 0.4465],  
                                                         [0.2023, 0.1994, 0.2010]))
```

1.3.2. Network:


As shown in above network 13 conv layer and 3 fully connected linear layers were used.

After each conv layer and linear layer there was a layer of Batch Normalization except the last linear layer.

A dropout of 50% was used before first 2 linear layer to avoid overfitting.

Relu was used as activation throughout the network.

Crossentropy loss with SGD optimizer was used with 0.01 learning rate for 150 epoch and train batch of 32.



1.3.3. Ablation Study

submission_netid8.csv 2 days ago by VishalBatavia more layers, 150 epoches	0.64500
submission_netid7.csv 3 days ago by VishalBatavia More Conv layers	0.62200
submission_netid6.csv 3 days ago by VishalBatavia 100 epoch	0.50000
submission_netid5.csv 3 days ago by VishalBatavia 50 epoches with SGD	0.50500
submission_netid4.csv 3 days ago by VishalBatavia SGD optimizer	0.39800
submission_netid3.csv 3 days ago by VishalBatavia Augmentation with 15 epoches	0.35500
submission_netid2.csv 4 days ago by VishalBatavia Maxpool, Batch size	0.34600
submission_netid1.csv 4 days ago by VishalBatavia Changes in Network layers	0.28699
submission_netid.csv 4 days ago by VishalBatavia Vanila Submission	0.21300

As shown from the above submissions with 7 conv layer, 2 linear layer and 2 dropout with Batch normalization after each conc and linear layer and with Relu activation the max test accuracy obtained was 50%.

In between Vanilla and the network with 50% accuracy as described above, various other options were used like different epochs, learning rate, momentum, Adam optimizer, different percentage of dropout as well as no dropout.

2. Part 2 Transfer learning

- Epoch = 50
- Batch = 32
- LR = 0.001
- Pretrained ResNet18 with FC layer replaced as shown below

```
self.fc_layer = nn.Sequential(  
    # nn.Dropout(),  
    nn.Linear(num_feats, 512),  
    nn.BatchNorm1d(512),  
    nn.ReLU(inplace=True),  
    # nn.Dropout(),  
    # nn.Linear(512, 512),  
    # nn.BatchNorm1d(512),  
    # nn.ReLU(inplace=True),  
    nn.Linear(512, 200)  
)
```

- Data augmentation as shown below

```
data_transforms = {  
    'train': transforms.Compose([  
        transforms.Resize(256),  
        transforms.RandomRotation(45),  
        transforms.RandomResizedCrop(224),  
        transforms.RandomHorizontalFlip(),  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406],  
                               [0.229, 0.224, 0.225])  
    ])
```

- Optimizer SGD

2.1. ResNet as a fixed feature Extractor.


Training accuracy was 85% while testing accuracy was 63%

```
TRAINING Epoch 44/50 Loss 0.0280 Accuracy 0.8480
TRAINING Epoch 45/50 Loss 0.0296 Accuracy 0.8230
TRAINING Epoch 46/50 Loss 0.0287 Accuracy 0.8433
TRAINING Epoch 47/50 Loss 0.0279 Accuracy 0.8367
TRAINING Epoch 48/50 Loss 0.0268 Accuracy 0.8437
TRAINING Epoch 49/50 Loss 0.0266 Accuracy 0.8463
TRAINING Epoch 50/50 Loss 0.0257 Accuracy 0.8533
Finished Training
-----
```


Test Loss: 0.0456 Test Accuracy 0.6301

- Visualization

class: 047.American_Goldfinch predicted: 047.American_Goldfinch



class: 079.Belted_Kingfisher predicted: 079.Belted_Kingfisher



class: 099.Ovenbird predicted: 099.Ovenbird



class: 146.Forsters_Tern predicted: 111.Loggerhead_Shrike



class: 188.Pileated_Woodpecker predicted: 188.Pileated_Woodpecker



2.2. Finetuning the ResNet

Training accuracy was 83% while testing accuracy was 55%

```
TRAINING Epoch 44/50 Loss 0.0861 Accuracy 0.8403
TRAINING Epoch 45/50 Loss 0.0944 Accuracy 0.8223
TRAINING Epoch 46/50 Loss 0.0829 Accuracy 0.8470
TRAINING Epoch 47/50 Loss 0.0912 Accuracy 0.8330
TRAINING Epoch 48/50 Loss 0.0885 Accuracy 0.8343
TRAINING Epoch 49/50 Loss 0.0819 Accuracy 0.8457
TRAINING Epoch 50/50 Loss 0.0844 Accuracy 0.8380
Finished Training
-----
```

```
In [28]: test(model2, criterion)
```

```
Test Loss: 0.2730 Test Accuracy 0.5532
```

- Visualization

class: 133.White_throated_Sparrow predicted: 133.White_throated_Sparrow



class: 020.Yellow_breasted_Chats predicted: 020.Yellow_breasted_Chats



class: 163.Cape_May_Warbler predicted: 183.Northern_Waterthrush



class: 046.Gadwall predicted: 052.Pied_billed_Grebe



class: 161.Blue_winged_Warbler predicted: 161.Blue_winged_Warbler

